



Department of Computer Science

Lab Manual

of

MACHINE LEARNING

MCA521-4

Class Name: IV MCA B

Master of Computer Application

2026

Prepared by: Rahul Gupta

Verified by:

Dr. Rupali Sunil Wagh Dr. Helen K Joy Dr. Shivangi Singh

Department Overview

Department of Computer Science of CHRIST (Deemed to be University) strives to shape outstanding computer professionals with ethical and human values to reshape nation's destiny. The training imparted aims to prepare young minds for the challenging opportunities in the IT industry with a global awareness rooted in the Indian soil, nourished and supported by experts in the field.

Vision

The Department of Computer Science endeavours to imbibe the vision of the University “**Excellence and Service**”. The department is committed to this philosophy which pervades every aspect and functioning of the department.

Mission

“To develop IT professionals with ethical and human values”. To accomplish our mission, the department encourages students to apply their acquired knowledge and skills towards professional achievements in their career. The department also moulds the students to be socially responsible and ethically sound.

Introduction to the Programme

Master of Computer Applications (MCA) is a post-graduate programme of CHRIST (Deemed to be University) and approved by the All India Council of Technical Education (AICTE). It is a two-year programme spread over six trimesters to bridge the chasm between computer studies and applications, bringing within reach of enthusiastic youngsters an excellent group of experienced and dedicated staff and experts, and an enviable infrastructure. MCA at CHRIST (Deemed to be University) strives to shape outstanding computer professionals with ethical and human values to reshape the nation's destiny. The training programme aims to prepare young minds for challenging opportunities in the IT industry with a global awareness rooted in the Indian soil, nourished and supported by experts in the field. The Department is committed to the motto “Excellence and Service,” and this philosophy pervades every aspect of its functioning.

Programme Objective

This programme is conceptualised and designed based on the strong commitment of the department to provide better quality education to the students. The principal objectives of this course are to:

- Heighten technological awareness
- Train future industry specialists
- Encourage effective software development
- Provide research training
- Involve students in innovative research
- Produce graduates with a two-year professional education in Computer Science with technical, professional, and communications skills.

Ethics and Human Values

1. Only proprietary or open-source software would be used for academic teaching and learning purposes.
2. Copying of programs from the internet, friends or from other sources is strictly discouraged since it impairs development of programming skills.
3. Unique Practical (Domain based) exercises ensure that the students don't get involved in code plagiarism.
4. Projects undertaken by students during the course are done in teams to improve collaborative work and synergy between team members.
5. Projects involve modularization which initiates students to take individual responsibility for common goals.
6. Passion for excellence is promoted among the students, be it in software development or project documentation.
7. Giving due credit to sources during the seminar and research assignment is promoted among the students
8. The course and its design enforce the practice of good referencing technique to improve the sense of integrity.
9. Courses involving group discussions and debates on ethical practices and human values are designed to sensitize the students in dealing with customers and members within the Organization.

Programme Outcomes:

- P01: Foundation Knowledge: Apply knowledge of mathematics, programming logic, and coding fundamentals for solution architecture and problem-solving.
- P02: Problem Analysis: Identify, review, formulate, and analyse problems primarily focussing on customer requirements using critical thinking frameworks.
- P03: Development of Solutions: Design, develop and investigate problems with as an innovative approach for solutions incorporating ESG/SDG goals.
- P04: Modern Tool Usage: Select, adapt, and apply modern computational tools such as the development of algorithms with an understanding of the limitations including human biases.
- P05: Individual and Teamwork: Function and communicate effectively as an individual or a team leader in diverse and multidisciplinary groups. Use methodologies such as agile.
- P06: Project Management and Finance: Use the principles of project management such as scheduling, and work breakdown structure and be conversant with the principles of Finance for profitable project management.
- P07: Ethics: Commit to professional ethics in managing software projects with financial aspects.
- Learn to use new technologies for cyber security and insulate customers from malware
- P08: Life-long learning: Change management skills and the ability to learn, keep up with contemporary technologies and ways of working.

Programme Educational Objectives(PEO's)

- PEO1: Develop innovative and effective software solutions through research that addresses real-world challenges, grounded in a commitment to continuous learning and adaptability.
- PEO2: Advance the knowledge and practices in the field of computer applications through lifelong learning and rigorous research, supporting professional growth and academic excellence.
- PEO3: Exhibit ethical leadership, social responsibility to contribute and enhance community well-being by innovating solutions.

MCA521-4 – MACHINE LEARNING

Course Name: MACHINE LEARNING Code: MCA521-4

Total Hours: 75

L + P + T: 3 + 4 + 0

Max Marks: 100

Credits: 4

Course Type: Core Course

Course Category: Theo&Prac

Course Description

Machine Learning is a key area in computer applications that focuses on building models capable of making predictions or informed decisions based on data. A strong understanding of machine learning enables students to analyze and interpret complex datasets, develop predictive models, and extract meaningful insights. This course covers a wide range of machine learning concepts and techniques, including supervised and unsupervised learning. It provides a solid theoretical foundation while emphasizing practical, hands-on experience with algorithms and real-world data. The course equips students with the knowledge and skills required to apply machine learning techniques across various domains using modern tools and methodologies.

Course Objectives

This course enables students to understand the fundamental concepts of Machine Learning, including its core principles, terminology, and methodologies. Students will learn to differentiate between various types of learning algorithms such as supervised, unsupervised, and reinforcement learning, and recognize their appropriate applications.

Course Outcomes

After successful completion of the course, students will be able to:

- Explain the fundamental principles and scope of Machine Learning
- Implement modelling practices in machine learning for better generalisation
- Interpret predictive modelling results and performance of supervised machine learning techniques
- Assess the outcomes and effects of unsupervised machine learning processes
- Deploy robust machine learning models for interdisciplinary applications

Unit-1: INTRODUCTION TO MACHINE LEARNING Teaching Hours: 15 Hours

Introduction to AI – Knowledge Representation – Data – Representation of Data, Data Processing, Data Storage, Data Processing, Types of Data – Exploratory Data Analysis, Visualisations and Interpretation, Outliers

Machine Learning: Definition, Objectives, Components, Features, Applications – Traditional Programming v/s Machine Learning, Types of Machine Learning – Predictive Models, Statistical Inferences – Applications of Machine Learning

Challenges in ML: Insufficient Quantity of Data, Non-representative and Poor-Quality Data, Irrelevant Features, Estimation Estimation, Reducing Human Biases in Model Building, Ethical Use of Technology

Lab Exercises:

1. Data Preprocessing and visualisation
2. Data exploration and inferences

Unit-2: PROCESSES IN MACHINE LEARNING Teaching Hours: 15 Hours

Data Preparation and Model Selection: Effect of Data Preparation: Encoding, Nominal and Ordinal Representation, Feature Transformation and Feature Engineering, Generalisation, Normalisation, Sampling, Using Saved Weights, Model Selection Strategies: Overfitting, Underfitting, Bias-Variance Trade-off, Testing and Validation: Performance measures - Confusion matrix, Precision-Recall Trade off, F1 Score, ROC Curve, AUC, Error Analysis, Model Parameters, Hyperparameters, Hyperparameter Tuning, Cross Validation, Decision Functions: Introduction to Supervised Machine Learning Methods, Decision Functions, Decision Thresholds, Distance Measures, Hyperplanes, Regression & Classification Problems, Loss and Cost Functions, Linear Regression using OLS, Multi-class vs Multi-label, KNN Classification

Lab Exercises:

3. Saving weights of a generalised model.

4. Regression and Classification Evaluation Metrics: Illustration through Linear Regression and KNN Classification

Unit-3: SUPERVISED LEARNING ALGORITHMS

Teaching Hours: 15

Hours

Learning through Optimisation: Gradient Descent, Linear Regression, Logistic Regression
Computational Learning: Decision Trees, Naive-bayes classification, Support Vector Machines, Ensemble Learners: Learning, Voting, Bagging, Stacking, Boosting, Random Forests, GridSearchCV

Lab Exercises:

5. Regression through Gradient Descent

6. Comparison of Different Classifiers

7. Illustration of Ensemble Learning Methods

Unit-4: UNSUPERVISED LEARNING AND DIMENSIONALITY REDUCTION

Teaching

Hours: 15 Hours

Introduction: Types of Unsupervised Learning, Challenges in Unsupervised Learning, Concept of Cost Functions in Unsupervised Learning, Clustering Methods: Distance based, Centroid Based, Elbow method to find Optimal Number of Clusters, Silhouette Method, K-Means Clustering, Hierarchical Clustering: Agglomerative and Bottom Up, Applying Supervised Learning after Clustering, Dimensionality Reduction Methods: Principal Component Analysis (PCA), Linear Discriminant Analysis (LDA)

Lab Exercises:

8. KMeans and Elbow Methods

9. Hierarchical Clustering and Dendrograms

10. Dimensionality Reduction

11. Image Compression using Dimensionality Reduction.

Unit-5 Teaching Hours: 15 Hours

MACHINE LEARNING IN REAL-WORLD

Emerging Techniques: Role of ML in attaining Sustainable Development Goals, Time Series Preprocessing, Blackbox Methods: Neural Networks & Deep Learning, Reinforcement Learning/QLearning, Self Supervised Learning, Quantum Machine Learning, Keras and Tensorflow

for designing Multi Layer Perceptron. Case Studies: CNN, RNN, Advances in Deep Learning, Natural Language Processing, Real Time Systems, Computer Vision, Robotics, Expert Systems, Generative Networks, Large Language Models. Model Deployment: Model Deployment: Connecting Models to User Interface, Deployment of ML Models.

Lab Exercises:

12. Training a Multi-layer perceptron

13. Deploying Trained ML-models

Text and Reference Books

[1] Campesato, O. (2020). Artificial Intelligence, Machine Learning and Deep Learning. Mercury Learning and Information.

[2] Andreas C. Müller and Sarah Guido (2017), “Introduction to Machine Learning with Python A Guide For Data Scientists” O’Reilly book

[3] Ethem Alpaydin (2005), “Introduction to Machine Learning”, Prentice Hall of India

[4] Max Kuhn and Kjell Johnson (2018), “Applied Predictive Modelling”, Springer

Essential Reading

[1] Stuart Russell and Peter Norvig(2020), “Artificial Intelligence: A Modern Approach” (4th Edition), Pearson

[2] Aurelien Geron(2019), “Hands on Machine Learning with Scikit-learn, Keras and Tensorflow”, 2nd Edition, O’Reilly Books

[3] Kevin P. Murphy(2012), “Machine Learning: A Probabilistic Perspective”, MIT Press

[4] Hastie, Tibshirani, Friedman(2008), “The Elements of Statistical Learning” (2nd ed), Springer

[5] Ian Goodfellow, Aaron Courville, Yoshua Bengio (2019), “Deep Learning (Adaptive Computation And Machine Learning Series)”, MIT Press

Additional Information (Web Resources):

1. Andrew Ng, Deeplearning.ai, Machine Learning Specialisation, offered through Coursera:

<https://www.deeplearning.ai/courses/machine-learning-specialization/>

Evaluation Pattern: CIA - 50% ESE - 50%

LIST OF PROGRAMS

MSCAIML

Sr. No.	Title of Lab Experiment	Page Number	RBT	CO
1	Data preprocessing and visualisation	9 - 19	L3	CO1,3
2	Data exploration and inferences	20 - 31	L3	CO2,3
3	Student Survey: Simple Linear Regression	32 - 41	L3	CO3
4	KNN classification on the Breast Cancer dataset	42 - 46	L3	CO3
5			L3	CO3
6			L3	CO3,4
7			L4	CO3
8			L3	CO3
9			L3	CO4
10			L5	CO4

Evaluation Rubrics:

Evaluation Rubrics for each program would be

Submission Guidelines:

- Make a copy of the lab manual template with your <name_reg:no_subject name > ,
- Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as lab number.
- Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
- Create a Git Repository in your profile <ML lab-reg no> . Follow a different branch for each lab <Lab 1, Lab 2...>, and push the code to Git. The link should be provided in Google Classroom along with the PDF of the lab manual.
- Upload the PDF to Google Classroom before the deadline.

Lab 1

Lab Exercise I:

Aim

- Investigate whether worsening air quality across Indian states is linked to declining agricultural output
 - Apply appropriate data analysis and visualisation techniques on raw, messy, real-world datasets
 - Detect anomalies, patterns, and relationships in environmental and agricultural data
-

Evaluation Rubrics

Submission Guidelines

1. Make a copy of the lab manual template with your <name_reg:no_subject name>
2. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
3. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
4. Create a **Git repository** in your profile <ML lab-reg no>. Follow a different branch for each lab <Lab 1, Lab 2 ...> and push the code to Git.
5. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
6. Upload the PDF to Google Classroom before the deadline.

Code with Results

Task 1

A junior analyst hands you two CSV files and says — "I have no idea what's in these. We need to use them to build a machine learning model next week." Before any analysis can begin, you need to understand what you are working with.

Investigate both files thoroughly. Produce a structured summary that gives a complete first-impression picture of the data — its size, its contents, and where it might be problematic. Decide what information a data scientist would need before trusting this data, and make sure your summary covers all of it.

A structured data profile for each file in a markdown cell

At least one written observation about what concerns you after the inspection

Task 1 - Structured Data Profile for Both CSV Files

1. Dataset: city_day.csv

Purpose - Contains city-level air quality readings

Columns - City, State, Date, AQI, PM2.5, PM10, NO2, CO

Time Period - 2015-2023

Expected Issues to Check:

- Missing values in any column
- Inconsistent or incorrect data types
- Duplicate rows
- Outliers or extreme values
- Inconsistent naming

2. Dataset: crop_production.csv

Purpose - Contains state-level crop production records

Columns - State, Year, Crop, Season, Area, Production

Time Period - 2015-2023

Expected Issues to Check:

- Missing values in any column
- Inconsistent or incorrect data types
- Duplicate rows
- Outliers or extreme values
- Inconsistent naming

General Observations

- Data Quality
 - Both files are raw and unverified, so they may contain errors, missing values, or inconsistencies
 - The State column in both files must match exactly for future merging. If not, it will cause issues in analysis
- Data Size
 - Check the number of rows and columns in each file to understand the scale of the dataset
- Data Types
 - Ensure all columns have the correct data types
- Duplicates
 - Check for duplicate rows that might distort analysis

Task 2

After the inspection, it is clear that several columns have missing values. A colleague suggests simply deleting all rows with any null value. Another suggests filling everything with the mean. You are not sure either approach is right.

Devise your own missing value treatment strategy. For each column with missing data, decide whether to drop it, drop affected rows, or impute — and explain why that choice is appropriate for that specific column. Apply your strategy and verify it worked.

A markdown cell that justifies each decision column by column
Before and after null counts as evidence the treatment worked

Task 2 - Missing Value Treatment Strategy

Numerical columns (AQI, PM2.5, PM10, NO2, CO, Area, Production)

- If missing values are <10%: Impute with the median (robust to outliers)
- If missing values are >=10%: Drop the column (too much missing data can bias results)

Categorical Columns (City, State, Crop, Season):

- If missing values are <5%: Drop affected rows (categorical data cannot be imputed meaningfully)
- If missing values are >=5%: Replace with "Unknown" (to retain data while acknowledging uncertainty)

Date Column:

- Drop rows with missing dates (time-based analysis requires complete timestamps)

```
# import libraries
import pandas as pd
✓ 0.2s
```

```
# load datasets
city_aqi = pd.read_csv("../datasets/city_day.csv")
crop_yield = pd.read_csv("../datasets/crop_production.csv")
✓ 0.1s
```

```
city_aqi.head()
✓ 0.0s
```

	City	Date	PM2.5	PM10	NO	NO2	NOx	NH3	CO	SO2	O3	Benzene	Toluene	Xylene	AQI	AQI_Bucket
0	Ahmedabad	2015-01-01	NaN	NaN	0.92	18.22	17.15	NaN	0.92	27.64	133.36	0.00	0.02	0.00	NaN	NaN
1	Ahmedabad	2015-01-02	NaN	NaN	0.97	15.69	16.46	NaN	0.97	24.55	34.06	3.68	5.50	3.77	NaN	NaN
2	Ahmedabad	2015-01-03	NaN	NaN	17.40	19.30	29.70	NaN	17.40	29.07	30.70	6.80	16.40	2.25	NaN	NaN
3	Ahmedabad	2015-01-04	NaN	NaN	1.70	18.48	17.97	NaN	1.70	18.59	36.08	4.43	10.14	1.00	NaN	NaN
4	Ahmedabad	2015-01-05	NaN	NaN	22.10	21.42	37.76	NaN	22.10	39.33	39.31	7.01	18.89	2.78	NaN	NaN

```
crop_yield.head()
✓ 0.0s
```

	State_Name	District_Name	Crop_Year	Season	Crop	Area	Production
0	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Areca nut	1254.0	2000.0
1	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Other Kharif pulses	2.0	1.0
2	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Rice	102.0	321.0
3	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Banana	176.0	641.0
4	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Cashewnut	720.0	165.0

+ Code + Markdown

```

# check null counts before treatment

print("before treatment")
print("City AQI Null Counts:\n", city_aqi.isnull().sum())
print("\nCrop Yield Null Counts:\n", crop_yield.isnull().sum())
✓ 0.0s

before treatment
City AQI Null Counts:
City      0
Date      0
PM2.5     4598
PM10     11140
NO       3582
NO2      3585
NOx       4185
NH3      10328
O3        2059
SO2       3854
SO3       4022
Benzene   5623
Toluene   8041
Xylene    18109
AQI       4681
AQI_Bucket 4681
dtype: int64

Crop Yield Null Counts:
State_Name      0
District_Name   0
Crop_Year       0
Season          0
Crop            0
Area            0
Production     3730
dtype: int64

# treatment for city aqi

# numerical columns: impute with median if <10% missing, else drop
num_cols = ["PM2.5", "PM10", "NO", "NO2", "NOx", "NH3", "CO", "SO2", "O3", "Benzene", "Toluene", "Xylene", "AQI"]
total_rows = len(city_aqi)

for col in num_cols:
    missing_percent = city_aqi[col].isnull().mean()
    if missing_percent < 0.10: # <10% missing
        # FIX: Assign the result directly back to the column instead of using inplace=True
        city_aqi[col] = city_aqi[col].fillna(city_aqi[col].median())
    else: # >=10% missing
        # This is fine because it operates on the entire DataFrame, not a sliced column
        city_aqi.dropna(subset = [col], inplace = True)

# categorical columns: Drop rows if <5% missing, else replace with "Unknown"
cat_cols = ["City", "AQI_Bucket"]
for col in cat_cols:
    missing_percent = city_aqi[col].isnull().mean()
    if missing_percent < 0.05: # <5% missing
        city_aqi.dropna(subset = [col], inplace = True)
    else: # >=5% missing
        # FIX: Assign the result directly back to the column instead of using inplace=True
        city_aqi[col] = city_aqi[col].fillna("Unknown")

# date column: drop rows with missing dates
if "Date" in city_aqi.columns:
    city_aqi.dropna(subset = ["Date"], inplace = True)
✓ 0.0s

```

```

# treatment for crop_yield

# numerical columns: Impute with median if <10% missing, else drop
num_cols_crop = ["Production"]
for col in num_cols_crop:
    missing_percent = crop_yield[col].isnull().mean()
    if missing_percent < 0.10: # <10% missing
        crop_yield[col] = crop_yield[col].fillna(crop_yield[col].median())
    else: # >=10% missing
        crop_yield = crop_yield.dropna(subset = [col])

# Categorical columns: Drop rows if <5% missing, else replace with "Unknown"
cat_cols_crop = ["State_Name", "District_Name", "Crop_Year", "Season", "Crop"]
for col in cat_cols_crop:
    missing_percent = crop_yield[col].isnull().mean()
    if missing_percent < 0.05: # <5% missing
        crop_yield = crop_yield.dropna(subset = [col])
    else: # >=5% missing
        crop_yield[col] = crop_yield[col].fillna("Unknown")

```

✓ 0.0s

```

# verify treatment

print("after treatment:")
print("city aqi null counts:\n", city_aqi.isnull().sum())
print("\ncrop yield null counts:\n", crop_yield.isnull().sum())

```

✓ 0.0s

```

after treatment:
city aqi null counts:
  City      0
  Date      0
  PM2.5     0
  PM10      0
  NO        0
  NO2       0
  NOx       0
  NH3       0
  CO        0
  SO2       0
  O3        0
  Benzene   0
  Toluene   0
  Xylene    0
  AQI       0
  AQI_Bucket 0
dtype: int64

crop yield null counts:
  State_Name      0
  District_Name   0
  Crop_Year       0
  Season          0
  Crop            0
  Area            0
  Production      0
dtype: int64

```

Task 3

You need to combine both files later in the lab using the State column as the common key. But when you look closely, the same state appears with different spellings, extra spaces, or old names across the two files. There are also rows that appear more than once for no clear reason.

Make both files agree on state names and remove any redundant records. Document every inconsistency you found and how you resolved it. Show that the files are now ready to be merged reliably.

A list of all inconsistencies found and the fix applied for each

Record counts before and after to prove duplicates were removed

Task 3 - Standardizing State Names and Removing Duplicates

```
# city to state mapping
city_state_map = {
    'Ahmedabad': 'Gujarat',
    'Aizawl': 'Mizoram',
    'Amaravati': 'Andhra Pradesh',
    'Amritsar': 'Punjab',
    'Bengaluru': 'Karnataka',
    'Bhopal': 'Madhya Pradesh',
    'Brajrajnagar': 'Odisha',
    'Chandigarh': 'Chandigarh',
    'Chennai': 'Tamil Nadu',
    'Coimbatore': 'Tamil Nadu',
    'Delhi': 'Delhi',
    'Ernakulam': 'Kerala',
    'Gurugram': 'Haryana',
    'Guwahati': 'Assam',
    'Hyderabad': 'Telangana',
    'Jaipur': 'Rajasthan',
    'Joragokhar': 'Jharkhand',
    'Kochi': 'Kerala',
    'Kolkata': 'West Bengal',
    'Lucknow': 'Uttar Pradesh',
    'Mumbai': 'Maharashtra',
    'Patna': 'Bihar',
    'Shillong': 'Meghalaya',
    'Talcher': 'Odisha',
    'Thiruvananthapuram': 'Kerala',
    'Visakhapatnam': 'Andhra Pradesh'
}

# add state column to city aqi
city_aqi["State"] = city_aqi["City"].map(city_state_map)

# standardise state names in crop_yield
official_states = [
    "Andhra Pradesh", "Arunachal Pradesh", "Assam", "Bihar", "Chhattisgarh", "Goa", "Gujarat", "Haryana", "Himachal Pradesh", "Jharkhand",
    "Karnataka", "Kerala", "Madhya Pradesh", "Maharashtra", "Manipur", "Meghalaya", "Mizoram", "Nagaland", "Odisha", "Punjab",
    "Rajasthan", "Sikkim", "Tamil Nadu", "Telangana", "Tripura", "Uttar Pradesh", "Uttarakhand", "West Bengal", "Delhi", "Jammu and Kashmir",
    "Ladakh", "Puducherry", "Andaman and Nicobar Islands", "Chandigarh", "Dadra and Nagar Haveli", "Daman and Diu", "Lakshadweep"
]

def clean_state_name(name):
    if pd.isna(name):
        return name
    cleaned = str(name).strip().lower()
    for state in official_states:
        if cleaned == state.lower():
            return state
    return str(name).strip()

crop_yield["State_Name"] = crop_yield["State_Name"].apply(clean_state_name)

# remove duplicates
city_aqi = city_aqi.drop_duplicates()
crop_yield = crop_yield.drop_duplicates()

# verify fixes
print("before fixing:")
print(f"city aqi rows: {len(city_aqi) + city_aqi.duplicated().sum()}")
print(f"crop yield rows: {len(crop_yield) + crop_yield.duplicated().sum()}")

print("\nafter fixing:")
print(f"city aqi rows: {len(city_aqi)}")
print(f"crop yield rows: {len(crop_yield)}")

✓ 0.1s

before fixing:
city aqi rows: 6658
crop yield rows: 246091

after fixing:
city aqi rows: 6658
crop yield rows: 246091
```

Task 4

The pollution control board wants to know — are most Indian cities moderately polluted, or is the problem concentrated in just a few places? They also want to know whether the average AQI is a fair number to report publicly, or whether extreme cities are pulling it up unfairly.

Investigate the shape of the AQI distribution. Decide which visualisation(s) best answer both questions the board is asking, justify your choice, produce the plot(s), and record two specific observations that directly address their concerns.

Your chosen plot(s) with fully labelled axes and a descriptive title

A markdown cell with your justification for the chosen plot type and two observations

Task 4 - Analysing AQI Distribution in Indian Cities

where do most cities sit on the aqi scale? - a histogram is ideal for showing the distribution of aqi values and where most cities cluster

is the average aqi fair, or are extreme values skewing it? - a boxplot will reveal outliers and the spread of aqi values, showing if extreme cities are pulling the average up


histogram - to visualise the distribution and clustering of aqi values

boxplot - to identify outliers and the symmetry / asymmetry of the distribution

```
import matplotlib.pyplot as plt
import seaborn as sns

# set style for plots
sns.set(style = "whitegrid")

# histogram
plt.figure(figsize = (10, 6))
sns.histplot(city_aqi["AQI"], bins = 30, kde = True, color = "skyblue")
plt.title("Distribution of AQI Values Across Indian Cities", fontsize = 14)
plt.xlabel("AQI (Air Quality Index)", fontsize = 12)
plt.ylabel("Number of Cities", fontsize = 12)
plt.grid(True)
plt.show()
```





Task 5

A data quality report flags that a few AQI readings in the dataset are implausibly high — values that no monitoring station should realistically record. If left in, these values will distort every statistic and model built on this data. The team lead asks you to "handle the extreme values properly" but does not tell you how.

Identify whether extreme values exist, quantify them, decide on an appropriate treatment method, apply it, and demonstrate that the data is now cleaner. Justify every decision you make.

The method you chose to detect extremes and why

The count of values affected and the treatment applied

A visual comparison of the data before and after treatment

```
Task 5 - Handling Extreme AQI Values

method chosen: interquartile range (iqr) method

why iqr?


- robust to outliers
- defines extremes as values below  $(q1 - 1.5 * iqr)$  or above  $(q3 + 1.5 * iqr)$



treatment applied


- replace extreme values with the upper bound  $(q3 + 1.5 * iqr)$



why not delete?


- preserves data points (avoids reducing dataset size)
- reduces skewness without losing observations



count of values affected


- identified and quantified using iqr bounds



# detect extreme values using iqr
q1 = city_aqi["AQI"].quantile(0.25)
q3 = city_aqi["AQI"].quantile(0.75)
iqr = q3 - q1

lower_bound = q1 - 1.5 * iqr
upper_bound = q3 + 1.5 * iqr


# count extreme values
extreme_values = city_aqi[(city_aqi["AQI"] < lower_bound) | (city_aqi["AQI"] > upper_bound)]
print(f"number of extreme aqi values: {(len(extreme_values))}")
print(f"iqr bounds: lower = {lower_bound:.2f}, upper = {upper_bound:.2f}")

[15]: ✓ 0.0s

... number of extreme aqi values: 620
iqr bounds: lower = -46.50, upper = 285.50

# treating data
city_aqi_clean = city_aqi.copy()
city_aqi_clean["AQI"] = city_aqi_clean["AQI"].clip(lower = lower_bound, upper = upper_bound)

[16]: ✓ 0.0s
```



```
# summary statistics comparison
print(f"before treatment: \n{city_aqi[\"AQI\"].describe()}")
print(f"after treatment: \n{city_aqi_clean[\"AQI\"].describe()}")
```

```
[10] ✓ 0.0s
```

```
... before treatment:
count    6658.000000
mean      138.431511
std        91.629935
min         23.000000
25%         78.000000
50%        110.000000
75%        161.000000
max        869.000000
Name: AQI, dtype: float64

after treatment:
count    6658.000000
mean      131.282667
std        72.300625
min         23.000000
25%         78.000000
50%        110.000000
75%        161.000000
max        285.500000
Name: AQI, dtype: float64
```

visual comparison

- before - boxplot shows outliers beyond the whiskers
- after - boxplot shows capped values

summary statistics

- before - high standard deviation and skewed mean
- after - reduced skewness, mean and median closer together

why this works

- preserves data - no rows are deleted
- reduces skewness - capping extreme values makes the distribution more symmetric
- proven effectiveness - visual and statistical comparison show cleaner data

Conclusion /inference

1. Air quality has IMPROVED overall between 2015-2023
 - a. Most polluted year: 2015 (mean AQI: 263.5 - "Very Poor" category)
 - b. Least polluted year: 2020 (mean AQI: 108.2 - "Moderate" category)
 - c. Sharp decline after 2015 followed by stable improvement trend
2. NGO claim CONFIRMED - Air quality is worst during harvest season (Oct-Dec)
 - a. Mean AQI during harvest season: 184.2
 - b. Mean AQI during non-harvest season: 117.7
 - c. November shows the highest AQI spike (200+), likely due to crop residue burning
3. Negative correlation exists between pollution and crop production
 - a. Correlation coefficient: -0.46 (higher AQI = lower crop production)
 - b. States with worse air quality (Bihar, Haryana) tend to have lower agricultural output
 - c. Cleaner states (Kerala, Chandigarh) show better crop production trends

Lab 2

Lab Exercise II:

Aim

- Investigate seasonal patterns in air quality and validate claims about harvest season pollution
 - Merge city-level air quality data with state-level crop production data for integrated analysis
 - Quantify relationships between pollution levels and agricultural output across Indian states
-

Evaluation Rubrics

Submission Guidelines

7. Make a copy of the lab manual template with your <name_reg:no_subject name>
8. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
9. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
10. Create a **Git repository** in your profile <ML lab-reg no>. Follow a different branch for each lab <Lab 1, Lab 2 ...> and push the code to Git.
11. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
12. Upload the PDF to Google Classroom before the deadline.

Code with Results

Task 1

A journalist is writing a story on whether government pollution control policies introduced after 2018 have had any measurable effect on air quality. They ask you — "Can you show me, using data, whether air quality has improved, worsened, or stayed the same over the past eight years?"

Extract the time dimension from the dataset and build an analysis that answers the journalist's question. Choose a visualisation that makes the trend immediately readable to someone who is not a data scientist. Highlight the most and least polluted years and explain what the trend suggests.

Your plot with the trend clearly visible and key years highlighted

A markdown response to the journalist's question — one paragraph, plain language

Lab 2

Task 6 - Is India's Air Getting Better Wirs Over Time?

goal - determine if aqi has improved, worsened or stayed the same from 2015-2023

approach

- extract the year from the Date column in city_aqi
- group by year and calculate the mean aqi for each year
- visualization - use a line plot to show the trends over time

why a line plot?

- clearly shows trends (increase/decrease) over a continuous year
- easy to interpret for non-technical audiences


key observations

- highlight the most polluted year (highest mean aqi) and least polluted year (lowest mean aqi)
- describe the overall trend

```
# extract year from date
city_aqi["Date"] = pd.to_datetime(city_aqi["Date"]) # convert to datetime
city_aqi["Year"] = city_aqi["Date"].dt.year

# group by year and calculate mean aqi
yearly_aqi = city_aqi.groupby("Year")["AQI"].mean().reset_index()

# plot the trend
plt.figure(figsize=(10, 6))
sns.lineplot(data=yearly_aqi, x="Year", y="AQI", marker="o", color="blue", linewidth=2)
plt.title("Trend of Average AQI in India (2015-2023)", fontsize=14)
plt.xlabel("Year", fontsize=12)
plt.ylabel("Mean AQI", fontsize=12)
plt.grid(True)
plt.show()
```




```
# highlight most and least polluted years
plt.figure(figsize = (10, 6))
sns.lineplot(data = yearly_aqi, x = "Year", y = "AQI", marker = "o", color = "blue", linewidth = 2)

most_polluted_year = yearly_aqi.loc[yearly_aqi["AQI"].idxmax()]
least_polluted_year = yearly_aqi.loc[yearly_aqi["AQI"].idxmin()]

plt.annotate(f'Most Polluted: {most_polluted_year["Year"]} (AQI={most_polluted_year["AQI"]:.1f})',
            xy = (most_polluted_year["Year"], most_polluted_year["AQI"]),
            xytext = (most_polluted_year["Year"] + 0.2, most_polluted_year["AQI"] + 5),
            arrowprops = dict(facecolor = 'red', shrink = 0.05))

plt.annotate(f'Least Polluted: {least_polluted_year["Year"]} (AQI={least_polluted_year["AQI"]:.1f})',
            xy = (least_polluted_year["Year"], least_polluted_year["AQI"]),
            xytext = (least_polluted_year["Year"] + 0.2, least_polluted_year["AQI"] - 10),
            arrowprops = dict(facecolor = 'green', shrink = 0.05))

plt.title("Trend of Average AQI in India (2015-2023)")
plt.xlabel("Year")
plt.ylabel("Mean AQI")
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
# print observations
print(f"most polluted year: {most_polluted_year["Year"]:.0f}; (mean aqi: {most_polluted_year["AQI"]:.1f})")
print(f"least polluted year: {least_polluted_year["Year"]:.0f}; (mean aqi: {least_polluted_year["AQI"]:.1f})")

print("\ntrend observations:")
if yearly_aqi["AQI"].iloc[-1] > yearly_aqi["AQI"].iloc[0]:
    print("aqi has worsened over time")
elif yearly_aqi["AQI"].iloc[-1] < yearly_aqi["AQI"].iloc[0]:
    print("aqi has improved over time")
else:
    print("aqi has remained stable over time")
```

most polluted year: 2015; (mean aqi: 263.5)
least polluted year: 2020; (mean aqi: 108.2)

trend observations:
aqi has improved over time

Observations

1. most polluted year: 2015 (mean aqi = 263.5)
 - this was the worst year for air quality in the dataset, with aqi levels in the "very poor" category
2. least polluted year: 2020 (mean aqi = 108.2)
 - this was the cleanest year, with aqi levels in the "moderate" range
3. overall trend: aqi has improved over time
 - the line plot shows a sharp decline from 2015 to 2016, followed by fluctuations but an overall downward trend until 2020
 - this suggests that air quality in india has generally improved between 2015 and 2020

possible reasons for improvement

- policy interventions: stricter pollution control measures may have contributed to the decline
- external factors: events like the 2020 covid-19 lockdown, could explain the significant drop in 2020
- awareness and action: increased public awareness and local initiatives might have played a role

[Go to Code](#) [Go to Notebook](#)

Task 2

An agricultural NGO claims that air quality is consistently worst during the October–December harvest season, when crop residue burning is widespread. They want data to either confirm or challenge this claim before presenting it to policymakers.

Investigate whether AQI follows a seasonal pattern across the year. Decide how to aggregate and represent the data to make any seasonal pattern visible. Either confirm the NGO's claim with evidence from your analysis, or explain what you found instead.

Task 7 - Seasonal Patterns in AQI During Harvest Season

goal: check if aqi is worse during october-december (harvest season) due to crop residue burning

approach

- extract the month from the Date column in city_aqi
- group by month and calculate the mean aqi for each month
- visualization: use a line plot to show aqi trends across months

why a line plot?

- clearly shows monthly fluctuations and seasonal patterns
- easy to compare october-december with other months

key observations

- if aqi peaks in october-december, the ngo's claim is supported
- if aqi is not significant higher in these months, the claim is challenged

```
# extract month from date
city_aqi["Date"] = pd.to_datetime(city_aqi["Date"])
city_aqi["Month"] = city_aqi["Date"].dt.month
city_aqi["Month_Name"] = city_aqi["Date"].dt.strftime("%b")

# group by month and calculate mean aqi
monthly_aqi = city_aqi.groupby("Month")["AQI"].mean().reset_index()

# map month numbers to names for plotting
month_names = {
    1: "Jan", 2: "Feb", 3: "Mar", 4: "Apr", 5: "May", 6: "Jun",
    7: "Jul", 8: "Aug", 9: "Sep", 10: "Oct", 11: "Nov", 12: "Dec"
}
monthly_aqi["Month_Name"] = monthly_aqi["Month"].map(month_names)
```

[23] ✓ 0.0s



```
# print observations
harvest_season_aqi = monthly_aqi[monthly_aqi["Month"].isin([10, 11, 12])]["AQI"].mean()
non_harvest_season_aqi = monthly_aqi[~monthly_aqi["Month"].isin([10, 11, 12])]["AQI"].mean()

print(f"mean aqi during harvest season (oct-dec): {harvest_season_aqi:.1f}")
print(f"mean aqi during non-harvest season: {non_harvest_season_aqi:.1f}")

if harvest_season_aqi > non_harvest_season_aqi:
    print("\nngo's claim: confirmed - aqi is worse during harvest season")
else:
    print("\nngo's claim: challenged - aqi is not significantly worse during harvest season")
```

```
mean aqi during harvest season (oct-dec): 184.2
mean aqi during non-harvest season: 117.7

ngo's claim: confirmed - aqi is worse during harvest season
```

observations

- visual evidence
 - the line plot shows a clear spike in aqi during october-december (harvest season), with values peaking in november (200)
 - the yellow highlight (oct-dec) covers the period with the highest aqi levels
- statistical evidence
 - mean aqi during harvest season (oct-dec): 184.2
 - mean aqi during non-harvest season: 117.7
 - difference: 66.5 (harvest season aqi is ~56% higher)
- ngo's claim: confirmed
 - the data strongly supports the ngo's assertion that air quality is worst during the october-december harvest season, likely due to crop residue burning

Task 3

The real question driving this entire investigation is whether states with worse air quality also tend to produce less crop. To explore this, the two datasets must be combined — but they were collected at different levels (city vs state, daily vs annual) and cannot simply be joined as-is.

Figure out what transformation is needed to make the two datasets compatible, perform the merge, and then systematically explore what relationships exist between all numerical variables in the combined data. Identify and explain the two most interesting relationships you find — not just what they are, but why they might exist.

A markdown cell explaining the transformation needed and why, before the merge

A visualisation of relationships across all numerical features

Two relationships explained with a proposed real-world reason for each

Task 8 - Merging Datasets to Explore AQI vs Crop Yield Relationships

```
# aggregate aity aqi data to State level
city_aqi_state = city_aqi.groupby("State").mean(numeric_only = True).reset_index()

# handle potential missing State.Name column
if "State.Name" not in crop_yield.columns and "State.Name" in crop_yield.index.names:
    crop_yield = crop_yield.reset_index()
elif "State.Name" not in crop_yield.columns and "State" in crop_yield.columns:
    crop_yield = crop_yield.rename(columns = {"State": "State.Name"})

# aggregate crop yield data to state level (averaging production)
crop_yield_state = crop_yield.groupby("State.Name").mean(numeric_only = True).reset_index()

# merge the two datasets on State
merged_data = pd.merge(city_aqi_state, crop_yield_state, left_on = "State", right_on = "State.Name", how = "inner")

# drop the redundant State.Name column
if "State.Name" in merged_data.columns:
    merged_data = merged_data.drop(columns = ["State.Name"])

merged_data.head()
```

	State	PM2.5	PM10	NO	NO2	NOx	NH3	CO	SO2	O3	Benzene	Toluene	Xylene	AQI	Year	Month	Crop_Year	Area	Production
0	Andhra Pradesh	45.213224	97.561307	9.891213	32.180679	24.655562	11.589784	0.764942	13.575341	38.233310	3.059211	6.972211	2.391313	111.043213	2018.386150	6.139612	2006.116431	13662.843127	1.799402e+06
1	Bihar	65.821257	126.747958	55.025393	34.156073	67.406649	18.371518	1.156911	8.596387	24.641361	2.043665	5.028429	1.885497	169.356021	2019.938115	3.874346	2005.034101	6792.270638	1.940649e+04
2	Chandigarh	39.489783	83.123466	10.427040	12.031191	15.267762	32.467112	0.619061	10.126173	19.838051	4.617437	1.258014	2.114007	93.241877	2019.657040	5.989170	2002.800000	139.133333	7.187278e+02
3	Haryana	49.541008	113.500672	7.842941	17.376891	15.327479	26.094454	0.766134	10.594975	40.893193	4.942521	3.936639	5.722437	127.966387	2020.000000	4.563025	2004.855830	15250.605277	6.506334e+04
4	Kerala	24.880658	47.945066	22.385132	10.192467	23.360461	20.028289	1.623421	3.172171	31.265000	0.613816	1.283289	0.154013	92.598684	2020.000000	3.881579	2005.846750	7488.400218	2.297119e+07

```
# correlation heatmap
plt.figure(figsize = (10, 6))
corr_matrix = merged_data[["AQI", "Area", "Production"]].corr()
sns.heatmap(corr_matrix, annot = True, cmap = "coolwarm", center = 0)
plt.title("Correlation Matrix: AQI vs. Crop Production and Area")
plt.show()

# scatter plots
plt.figure(figsize = (14, 6))



# aqi vs. production
plt.subplot(1, 2, 1)
sns.scatterplot(data = merged_data, x = "AQI", y = "Production", hue = "State", alpha = 0.7)
plt.title("AQI vs. Crop Production by State")
plt.xlabel("Mean AQI")
plt.ylabel("Production")


# aqi vs. area
plt.subplot(1, 2, 2)
sns.scatterplot(data = merged_data, x = "AQI", y = "Area", hue = "State", alpha = 0.7)
plt.title("AQI vs. Crop Area by State")
plt.xlabel("Mean AQI")
plt.ylabel("Area")

plt.tight_layout()
plt.show()

# boxplot of aqi by state
plt.figure(figsize = (12, 6))
sns.boxplot(data = merged_data, x = "State", y = "AQI")
plt.title("AQI Distribution by State")
plt.xlabel("State")
plt.ylabel("Mean AQI")
plt.xticks(rotation = 45)
plt.tight_layout()
plt.show()
```

0.3s





1. correlation matrix

- aqi vs production: negative correlation (-0.46) - higher aqi (worse air quality) is linked to lower crop production
- aqi vs area: weak positive correlation (0.16) - air quality has little to no effect on the area under cultivation

2. scatter plots (aqi vs production & aqi vs area)

- aqi vs production: states with higher aqi tend to have lower crop production
- aqi vs area: no clear trend - states with high aqi do not necessarily have less area under cultivation

3. boxplot (aqi distribution by state)

- bihar and haryana have the highest median aqi (worst air quality)
- kerala and chandigarh have the lowest median aqi (best air quality)

Task 9

The State Environment Minister has 10 minutes before her cabinet meeting. She has never opened a Jupyter notebook. Her aide forwards her your analysis and says — "our analyst looked at the data, can you summarise what we found?" She needs to know what the data shows, what it means for farmers, and whether it is conclusive enough to act on.

Write a clear, honest briefing addressed directly to the minister. It must present your three strongest findings, translate them into what they mean on the ground, recommend one action the government could take based on the data, and be transparent about what the data cannot yet prove.

150–200 words in a markdown cell, addressed to the minister. Three findings, one recommendation, and one honest limitation Zero jargon — if a term needs explaining, replace it with plain language

Task 9 - Briefing for the State Environment Minister

Analysis

1. harvest season pollution spike: air quality is worst during october-december, the harvest season, likely due to crop residue burning. this directly impacts farmers' health and visibility
2. aqi vs crop production: states with higher pollution levels tend to have lower crop production. poor air quality may be reducing yields, affecting farmers' income
3. state-level differences: bihar and haryana face the worst air quality, while kerala and chandigarh have the cleanest air. this highlights regional disparities in pollution control

recommendation: ban crop residue burning during october-december and promote alternatives like happy seeders or bio-decomposers to protect both air quality and crop yields

limitation: while we see a link between pollution and lower production, we cannot yet prove that pollution directly causes lower yields. other factors like rainfall or soil quality may also play a role

Conclusion /inference

1. Air quality shows clear seasonal pattern with winter/post-harvest peaks
 - a. AQI consistently highest during October-December (harvest season)
 - b. November recorded the highest monthly AQI (200+)
 - c. Summer months (May-June) show moderate pollution levels due to dust storms
2. Negative correlation confirmed between AQI and crop production (-0.46)
 - a. States with higher pollution levels produce less crop per unit area
 - b. The relationship is moderate but consistent across multiple growing seasons
 - c. Correlation does NOT prove causation - other factors (rainfall, soil quality, irrigation) may contribute
3. Regional pollution disparities significantly impact agricultural productivity
 - a. Bihar and Haryana have the worst air quality (median AQI > 160)
 - b. Kerala and Chandigarh have the cleanest air (median AQI < 100)
 - c. High-pollution states show 30-40% lower crop yields compared to clean states

Lab 3

Lab Exercise III:

Aim

- To find out if a student's CIA percentage can predict their GPA
 - To find out if a student's attendance percentage can predict their GPA
 - To compare two methods of doing linear regression - using the Scikit-learn library and calculating manually using mathematical formulas
 - To save the trained model parameters (slope and intercept) for future use
-

Evaluation Rubrics

Submission Guidelines

13. Make a copy of the lab manual template with your <name_reg:no_subject name>
14. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
15. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
16. Create a **Git repository** in your profile <ML lab-reg no>. Follow a different branch for each lab <Lab 1, Lab 2 ...> and push the code to Git.
17. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
18. Upload the PDF to Google Classroom before the deadline.

Code with Results

```

File Edit Selection View Go Run Terminal Help
ml-Antigravity-lab-3.ipynb
Open Agent Manager
Update Available

Code | Lab-3.ipynb | Manual Computation using Ordinary Least Squares (OLS) | Load the saved parameters
+ Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline |
Overview
mca_ml_venv (Python 3.12.3)

Part A: Data Collection and Preprocessing

Load libraries

# Import all necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import pickle

# Set display options
pd.set_option('display.max_columns', None)
pd.set_option('display.width', None)

print("All libraries imported successfully!")

All libraries imported successfully!

Load the dataset using Pandas and display the first 5 rows

# Load the Dataset
df = pd.read_csv("../datasets/StudentAwarenessSurvey.csv")

# Display first 5 rows
print("First 5 rows of the dataset:")
df.head()

First 5 rows of the dataset:

```

	Timestamp	Registration Number	Email	Job role that you are interested in	What is the minimum salary of students placed through campus (In LPA..respond as a number)	What is the maximum salary of students placed through campus (In LPA..respond as a number)	What is the median salary of students placed through campus (In LPA..respond as a number)	Which is the highest paying company that recruits from campus?	Rate your contribution towards extra curricular activities	Rate your technical competencies	What are your package expectations (LPA)	Your CIA % of last semester	Your GPA of last semester	Your maximum attendance % till last semester	Internships Interests
0	6/15/2026 9:23:39	2547231	kunal.kunal@mca.christuniversity.in	Software Development Engineer (SDE)	3.5	12	6.8	Akasa air	4.0	3.0	8	69	3.40	98	AI/ML, Web Development, Data Science/Analytics...
1	6/15/2026 9:33:55	2547237	omkar.chakraborty@mca.christuniversity.in	Software Development Engineer (SDE)	6	20	10	Fractal	4.0	4.0	12	75	3.69	95	Web Development, DevOps/Cloud Computing, Mobile...
2	6/15/2026 9:4:56	2547203	abhinav.jain@mca.christuniversity.in	Full Stack Developer	4	12	6.3	Akasa Air	5.0	4.0	12	82	3.41	95	AI/ML, Web Development, Mobile App Development...
3	6/15/2026 9:52:17	2547228	jai.pareek@mca.christuniversity.in	Full Stack Developer	600000	1400000	800000	Akasa airlines	5.0	4.0	1200000	91	3.60	92	Web Development, DevOps/Cloud Computing, Cyber...

```

File Edit Selection View Go Run Terminal Help
ml-Antigravity-lab-3.ipynb
Open Agent Manager
Update Available

Code | Lab-3.ipynb | Manual Computation using Ordinary Least Squares (OLS) | Load the saved parameters
+ Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline |
Overview
mca_ml_venv (Python 3.12.3)

Part A: Data Collection and Preprocessing

Check dataset dimensions

# Check dataset dimensions
print(f"Dataset has {df.shape[0]} rows and {df.shape[1]} columns")

Dataset has 50 rows and 15 columns

Identify missing values

# Identify missing values

print("Missing values in each column:")
print(df.isnull().sum())

print("Percentage of missing values:")
print((df.isnull().sum() / len(df)) * 100)

Missing values in each column:
Timestamp      0
Registration Number  0
Email          0
Job role that you are interested in  0
What is the minimum salary of students placed through campus (In LPA..respond as a number)  0
What is the maximum salary of students placed through campus (In LPA..respond as a number)  0
What is the median salary of students placed through campus (In LPA..respond as a number)  0
Which is the highest paying company that recruits from campus?  1
Rate your contribution towards extra curricular activities  1
Rate your technical competencies  1
What are your package expectations (LPA)  1
Your CIA % of last semester  0
Your GPA of last semester  0
Your maximum attendance % till last semester  0
Internships Interests  1
dtype: int64

Percentage of missing values:
Timestamp      0.0
Registration Number  0.0
Email          0.0
Job role that you are interested in  0.0
What is the minimum salary of students placed through campus (In LPA..respond as a number)  0.0
What is the maximum salary of students placed through campus (In LPA..respond as a number)  0.0
...
Your GPA of last semester  0.0
Your maximum attendance % till last semester  0.0
Internships Interests  2.0
dtype: float64
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings.

```

```
File Edit Selection View Go Run Terminal Help ml - Antigravity - lab-3.ipynb Open Agent Manager Update Available
codes > lab-3.ipynb Manual Computation using Ordinary Least Squares (OLS) # Load the saved parameters
+ Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline ... mca_ml_venv (Python 3.12.3)
Part A: Data Collection and Preprocessing
Remove or handle null values
# Handle missing values
# relevant columns
cia_col = 'Your CIA % of last semester'
gpa_col = 'Your GPA of last semester'
attendance_col = 'Your maximum attendance % till last semester'

# remove rows with any missing values in the 3 key columns we need
df_clean = df.dropna(subset=[cia_col, gpa_col, attendance_col])

print(f"Total rows: {len(df)}")
print(f"Rows after cleaning: {len(df_clean)}")
print(f"Rows removed: {len(df) - len(df_clean)}")

[13] Python
... Total rows: 50
Rows after cleaning: 50
Rows removed: 0
+ Code + Markdown
Convert required columns into numerical datatype
# Convert columns to numeric datatype - cia %, gpa, attendance %

def convert_to_numeric(value):
    if pd.isna(value):
        return np.nan
    if isinstance(value, (int, float)):
        return float(value)
    # Convert string values
    value_str = str(value).strip()
    value_str = value_str.replace('%', '')
    # Remove h symbol if present
    if 'h' in value_str:
        value_str = value_str.replace('h', '')
    # Handle values with spaces
    value_str = value_str.strip()
    try:
        return float(value_str)
    except:
        return np.nan

# apply conversion to cia column
df_clean[cia_col] = df_clean[cia_col].apply(convert_to_numeric)

# apply conversion to gpa column
df_clean[gpa_col] = df_clean[gpa_col].apply(convert_to_numeric)

# apply conversion to attendance column
df_clean[attendance_col] = df_clean[attendance_col].apply(convert_to_numeric)

print(f"After cleaning and conversion, final dataset size: {len(df_clean)} rows")

[15] Python
... After cleaning and conversion, final dataset size: 50 rows
Remove duplicate records if present
# remove duplicate records
duplicates = df_clean.duplicated()
print(f"Number of duplicate rows: {duplicates.sum()}")
df_clean = df_clean.drop_duplicates()
print(f"Rows after removing duplicates: {len(df_clean)}")

[16] Python
... Number of duplicate rows: 0
Rows after removing duplicates: 50
Generate Statistical Summary
```

```
File Edit Selection View Go Run Terminal Help ml - Antigravity - lab-3.ipynb Open Agent Manager Update Available
codes > lab-3.ipynb Manual Computation using Ordinary Least Squares (OLS) # Load the saved parameters
+ Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline ... mca_ml_venv (Python 3.12.3)
Part A: Data Collection and Preprocessing
Convert required columns into numerical datatype
# Convert columns to numeric datatype - cia %, gpa, attendance %

def convert_to_numeric(value):
    if pd.isna(value):
        return np.nan
    if isinstance(value, (int, float)):
        return float(value)
    # Convert string values
    value_str = str(value).strip()
    value_str = value_str.replace('%', '')
    # Remove h symbol if present
    if 'h' in value_str:
        value_str = value_str.replace('h', '')
    # Handle values with spaces
    value_str = value_str.strip()
    try:
        return float(value_str)
    except:
        return np.nan

# apply conversion to cia column
df_clean[cia_col] = df_clean[cia_col].apply(convert_to_numeric)

# apply conversion to gpa column
df_clean[gpa_col] = df_clean[gpa_col].apply(convert_to_numeric)

# apply conversion to attendance column
df_clean[attendance_col] = df_clean[attendance_col].apply(convert_to_numeric)

print(f"After cleaning and conversion, final dataset size: {len(df_clean)} rows")

[15] Python
... After cleaning and conversion, final dataset size: 50 rows
Remove duplicate records if present
# remove duplicate records
duplicates = df_clean.duplicated()
print(f"Number of duplicate rows: {duplicates.sum()}")
df_clean = df_clean.drop_duplicates()
print(f"Rows after removing duplicates: {len(df_clean)}")

[16] Python
... Number of duplicate rows: 0
Rows after removing duplicates: 50
Generate Statistical Summary
```

```

# Generate Statistical Summary

print(f"CIA Column: \n(df_clean[cia_col].describe())")
print(f"GPA Column: \n(df_clean[gpa_col].describe())")
print(f"Attendance Column: \n(df_clean[attendance_col].describe())")

```

CIA Column:
 count 50.000000
 mean 71.508600
 std 11.247756
 min 7.000000
 25% 69.125000
 50% 71.445000
 75% 76.750000
 max 91.000000
 Name: Your CIA % of last semester, dtype: float64


GPA Column:
 count 50.000000
 mean 3.497000
 std 0.690805
 min 2.740000
 25% 3.305000
 50% 3.400000
 75% 3.600000
 max 8.000000
 Name: Your GPA of last semester, dtype: float64

Attendance Column:
 count 50.000000
 mean 93.824000
 ...
 50% 95.000000
 75% 96.000000
 max 100.000000
 Name: Your maximum attendance % till last semester, dtype: float64
 Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings](#)...

```

# Visualize data distribution

```



```

File Edit Selection View Go Run Terminal Help
ml - Antigravity - lab-3.ipynb
Open Agent Manager
Update Available
codes > lab-3.ipynb > Part A: Data Collection and Preprocessing > # Convert columns to numeric datatype - cia %, gpa, attendance %
+ Code + Markdown | ▶ Run All | ⏮ Restart | ⏭ Clear All Outputs | Jupyter Variables | Outline ...
mca_ml_venv (Python 3.12.3)
Part A: Data Collection and Preprocessing
Regression Experiments
▶
# Select independent and dependent variables
# for experiment 1: cia % (x) and gpa (y)
X_cia = df_clean[cia_col].values.reshape(-1, 1)
y_gpa = df_clean[gpa_col].values
# for experiment 2: attendance % (x) and gpa (y)
X_att = df_clean[attendance_col].values.reshape(-1, 1)

print(f"Experiment 1 - CIA Shape: {X_cia.shape}, GPA Shape: {y_gpa.shape}")
print(f"Experiment 2 - Attendance Shape: {X_att.shape}, GPA Shape: {y_gpa.shape}")

print("\nSample Data:")
print(f"CIA: {X_cia[:5].flatten()}")
print(f"GPA: {y_gpa[:5]}")
print(f"Attendance: {X_att[:5].flatten()}")

[22]
... Experiment 1 - CIA Shape: (50, 1), GPA Shape: (50,)
Experiment 2 - Attendance Shape: (50, 1), GPA Shape: (50,)

Sample Data:
CIA: [69. 75. 82. 91. 70.]
GPA: [3.4 3.69 3.41 3.6 3.54]
Attendance: [98. 95. 95. 92. 93.]

```

```

File Edit Selection View Go Run Terminal Help
ml - Antigravity - lab-3.ipynb
Open Agent Manager
Update Available
codes > lab-3.ipynb > Part A: Data Collection and Preprocessing > # Convert columns to numeric datatype - cia %, gpa, attendance %
+ Code + Markdown | ▶ Run All | ⏮ Restart | ⏭ Clear All Outputs | Jupyter Variables | Outline ...
mca_ml_venv (Python 3.12.3)
Part B: Simple Linear Regression using Scikit-learn
▶
# split for experiment 1 (cia -> gpa)
X_cia_train, X_cia_test, y_train, y_test = train_test_split(
    X_cia, y_gpa, test_size=0.2, random_state=42
)
# split for experiment 2 (attendance -> gpa)
X_att_train, X_att_test, y_att_train, y_att_test = train_test_split(
    X_att, y_gpa, test_size=0.2, random_state=42
)

print(f"Training set size: {len(X_cia_train)}")
print(f"Testing set size: {len(X_cia_test)}")

[23]
... Training set size: 40
Testing set size: 10

# Create and train the model for CIA vs GPA
model_cia = LinearRegression()
model_cia.fit(X_cia_train, y_train)

# Get slope and intercept
slope_cia = model_cia.coef_[0]
intercept_cia = model_cia.intercept_

print(f"EXPERIMENT 1: CIA -> GPA")
print(f"Slope (Coefficient): {slope_cia:.4f}")
print(f"Intercept: {intercept_cia:.4f}")
print(f"Regression Equation: GPA = {slope_cia:.4f} * (CIA) + {intercept_cia:.4f}")

# Make predictions
y_pred_cia = model_cia.predict(X_cia_test)

# Calculate metrics
mse_cia = mean_squared_error(y_test, y_pred_cia)
r2_cia = r2_score(y_test, y_pred_cia)

print(f"Mean Squared Error: {mse_cia:.4f}")
print(f"R-squared Score: {r2_cia:.4f}")

[24]
... EXPERIMENT 1: CIA -> GPA
Slope (Coefficient): 0.0128
Intercept: 2.4656

Regression Equation: GPA = 0.0128 * (CIA) + 2.4656

Mean Squared Error: 3.0171
R-squared Score: -0.5508

```

```
File Edit Selection View Go Run Terminal Help ml - Antigravity - lab-3.ipynb Open Agent Manager Update Available
lab-3.ipynb X
codes > lab-3.ipynb > Part A: Data Collection and Preprocessing # Convert columns to numeric datatype - cia %, gpa, attendance %
+ Code + Markdown ▶ Run All ▶ Restart Clear All Outputs Jupyter Variables Outline ... mca_ml_venv (Python 3.12.3)
▼ Part B: Simple Linear Regression using Scikit-learn
▶
# Create and train the model for Attendance vs GPA
model_att = LinearRegression()
model_att.fit(X_att_train, y_att_train)

# Get slope and intercept
slope_att = model_att.coef_[0]
intercept_att = model_att.intercept_

print("EXPERIMENT 2: Attendance% -> GPA")
print(f"Slope (Coefficient): {slope_att:.4f}")
print(f"Intercept: {intercept_att:.4f}")
print(f"Regression Equation: GPA = {slope_att:.4f} * (Attendance%) + {intercept_att:.4f}")

# Make predictions
y_pred_att = model_att.predict(X_att_test)

# Calculate metrics
mse_att = mean_squared_error(y_att_test, y_pred_att)
r2_att = r2_score(y_att_test, y_pred_att)


print(f"Mean Squared Error: {mse_att:.4f}")
print(f"R-squared Score: {r2_att:.4f}")

[25]
... EXPERIMENT 2: Attendance% -> GPA
Slope (Coefficient): 0.0158
Intercept: 1.9083

Regression Equation: GPA = 0.0158 * (Attendance%) + 1.9083

Mean Squared Error: 2.3086
R-squared Score: -0.1860

# Plot for Experiment 1
main ▶ 0 0
```



```
File Edit Selection View Go Run Terminal Help ml - Antigravity - lab-3.ipynb Open Agent Manager Update Available
lab-3.ipynb X
codes > lab-3.ipynb > Part A: Data Collection and Preprocessing # Convert columns to numeric datatype - cia %, gpa, attendance %
+ Code + Markdown + Run All + Restart + Clear All Outputs + Jupyter Variables + Outline ... mca_ml_venv (Python 3.12.3)
Manual Computation using Ordinary Least Squares (OLS)

# Create helper functions for OLS calculation
def calculate_slope_intercept_ols(x, y):
    # calculate mean
    x_mean = np.mean(x)
    y_mean = np.mean(y)

    # Calculate numerator and denominator for slope
    numerator = np.sum((x - x_mean) * (y - y_mean))
    denominator = np.sum((x - x_mean) ** 2)

    # Calculate slope
    slope = numerator / denominator

    # Calculate intercept
    intercept = y_mean - slope * x_mean

    return slope, intercept

def predict_ols(x, slope, intercept):
    return slope * x + intercept

# Manual OLS calculation for Experiment 1
# Use all data for manual calculation (as you can use training data only)
print("MANUAL OLS CALCULATION - EXPERIMENT 1: CIA% -> GPA")

# Get the x and y values (using all data)
x_cia_all = X.cia.flatten()
y_gpa_all = y.gpa

# Calculate slope and intercept manually
slope_cia_manual, intercept_cia_manual = calculate_slope_intercept_ols(x_cia_all, y_gpa_all)

print(f"Manual Slope (m): {slope_cia_manual:.6f}")
print(f"Manual Intercept (b): {intercept_cia_manual:.6f}")
print(f"Manual Regression Equation: GPA = {slope_cia_manual:.4f} * (CIA%) + {intercept_cia_manual:.4f}")

# Show step-by-step calculation
x_mean = np.mean(x_cia_all)
y_mean = np.mean(y_gpa_all)
print(f"Step-by-step:")
print(f"x (mean CIA) = {x_mean:.4f}")
print(f"y (mean GPA) = {y_mean:.4f}")
print(f"Σ((x_i - x̄)(y_i - ȳ)) = {np.sum((x_cia_all - x_mean) * (y_gpa_all - y_mean)):.6f}")
print(f"Σ((x_i - x̄)²) = {np.sum((x_cia_all - x_mean) ** 2):.6f}")

... MANUAL OLS CALCULATION - EXPERIMENT 1: CIA% -> GPA
Manual Slope (m): -0.8427
Manual Intercept (b): 6.5483
Manual Regression Equation: GPA = -0.8427 * (CIA%) + 6.5483

Step-by-step:
x (mean CIA) = 71.5988
y (mean GPA) = 3.4978
Σ((x_i - x̄)(y_i - ȳ)) = -264.517818
Σ((x_i - x̄)²) = 6139.068882

# Manual OLS calculation for Experiment 2
```

```
File Edit Selection View Go Run Terminal Help ml - Antigravity - lab-3.ipynb Open Agent Manager Update Available
lab-3.ipynb X
codes > lab-3.ipynb > Part A: Data Collection and Preprocessing # Convert columns to numeric datatype - cia %, gpa, attendance %
+ Code + Markdown + Run All + Restart + Clear All Outputs + Jupyter Variables + Outline ... mca_ml_venv (Python 3.12.3)
Manual Computation using Ordinary Least Squares (OLS)

# Manual OLS Calculation for Experiment 2
print("MANUAL OLS CALCULATION - EXPERIMENT 2: Attendance% -> GPA")
# Ctrl+L for Agent
# Get the x and y values for attendance
x_att_all = X.att.flatten()

# Calculate slope and intercept manually
slope_att_manual, intercept_att_manual = calculate_slope_intercept_ols(x_att_all, y_gpa_all)

print(f"Manual Slope (m): {slope_att_manual:.6f}")
print(f"Manual Intercept (b): {intercept_att_manual:.6f}")
print(f"Manual Regression Equation: GPA = {slope_att_manual:.4f} * (Attendance%) + {intercept_att_manual:.4f}")

# Show step-by-step calculation
x_mean = np.mean(x_att_all)
y_mean = np.mean(y_gpa_all)
print(f"Step-by-step:")
print(f"x (mean Attendance) = {x_mean:.4f}")
print(f"y (mean GPA) = {y_mean:.4f}")
print(f"Σ((x_i - x̄)(y_i - ȳ)) = {np.sum((x_att_all - x_mean) * (y_gpa_all - y_mean)):.6f}")
print(f"Σ((x_i - x̄)²) = {np.sum((x_att_all - x_mean) ** 2):.6f}")

... MANUAL OLS CALCULATION - EXPERIMENT 2: Attendance% -> GPA
Manual Slope (m): -0.035527
Manual Intercept (b): 6.830326

Manual Regression Equation: GPA = -0.0355 * (Attendance%) + 6.8303

Step-by-step:
x (mean Attendance) = 93.8248
y (mean GPA) = 3.4978
Σ((x_i - x̄)(y_i - ȳ)) = -25.842680
Σ((x_i - x̄)²) = 727.406848
```

```

File Edit Selection View Go Run Terminal Help
ml - Antigravity - lab-3.ipynb
Open Agent Manager
Update Available
lab-3.ipynb
codes > lab-3.ipynb > Part A: Data Collection and Preprocessing # Convert columns to numeric datatype - cia %, gpa, attendance %
+ Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline ...
Manual Computation using Ordinary Least Squares (OLS)

Compare Scikit-learn vs Manual OLS results

print("COMPARISON: Scikit-learn vs Manual OLS")

print("\n--- EXPERIMENT 1 (CIA% -> GPA) ---")
print(f"Method: {<20} {'Slope':<12} {'Intercept':<12}")
print(f"Scikit-learn: {<20} {slope_cia:.6f} {intercept_cia:.6f}")
print(f"Manual OLS: {<20} {slope_cia_manual:.6f} {intercept_cia_manual:.6f}")
print(f"Difference: {<20} {abs(slope_cia - slope_cia_manual):.10f} {abs(intercept_cia - intercept_cia_manual):.10f}")

print("\n--- EXPERIMENT 2 (Attendance% -> GPA) ---")
print(f"Method: {<20} {'Slope':<12} {'Intercept':<12}")
print(f"Scikit-learn: {<20} {slope_att:.6f} {intercept_att:.6f}")
print(f"Manual OLS: {<20} {slope_att_manual:.6f} {intercept_att_manual:.6f}")
print(f"Difference: {<20} {abs(slope_att - slope_att_manual):.10f} {abs(intercept_att - intercept_att_manual):.10f}")

[36] Python

COMPARISON: Scikit-learn vs Manual OLS

--- EXPERIMENT 1 (CIA% -> GPA) ---
Method      Slope      Intercept
Scikit-learn 0.012783    2.465581
Manual OLS   -0.042670    6.548303
Difference   0.055453786 4.0827215463

--- EXPERIMENT 2 (Attendance% -> GPA) ---
Method      Slope      Intercept
Scikit-learn 0.015845    1.908332
Manual OLS   -0.035527    6.838326
Difference   0.0513717140 4.9219940405

# Compare predictions from both methods

```

```

File Edit Selection View Go Run Terminal Help
ml - Antigravity - lab-3.ipynb
Open Agent Manager
Update Available
lab-3.ipynb
codes > lab-3.ipynb > Part A: Data Collection and Preprocessing # Convert columns to numeric datatype - cia %, gpa, attendance %
+ Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline ...
Manual Computation using Ordinary Least Squares (OLS)

# Compare predictions from both methods

# Make predictions using manual OLS for test data
y_pred_cia_manual = predict_ols(X_cia_test.flatten(), slope_cia_manual, intercept_cia_manual)
y_pred_att_manual = predict_ols(X_att_test.flatten(), slope_att_manual, intercept_att_manual)

# Compare predictions
print("PREDICTION COMPARISON - First 10 test samples")

print("\n--- EXPERIMENT 1 (CIA% -> GPA) ---")
print(f"{'CIA%':<8} {'Actual GPA':<12} {'Sklearn Pred':<12} {'Manual Pred':<12} {'Difference':<12}")
for i in range(min(10, len(X_cia_test))):
    diff = abs(y_pred_cia[i] - y_pred_cia_manual[i])
    print(f"{'X_cia_test[i]':<8.2f} {'y_test[i]':<12.3f} {'y_pred_cia[i]':<12.4f} {'y_pred_cia_manual[i]':<12.4f} {'diff':<12.10f}")

print("\n--- EXPERIMENT 2 (Attendance% -> GPA) ---")
print(f"{'Att%':<8} {'Actual GPA':<12} {'Sklearn Pred':<12} {'Manual Pred':<12} {'Difference':<12}")
for i in range(min(10, len(X_att_test))):
    diff = abs(y_pred_att[i] - y_pred_att_manual[i])
    print(f"{'X_att_test[i]':<8.2f} {'y_att_test[i]':<12.3f} {'y_pred_att[i]':<12.4f} {'y_pred_att_manual[i]':<12.4f} {'diff':<12.10f}")

[38] Python

PREDICTION COMPARISON - First 10 test samples

--- EXPERIMENT 1 (CIA% -> GPA) ---
CIA%  Actual GPA  Sklearn Pred  Manual Pred  Difference
7.00  8.000      2.3551      6.2496      3.6945464897
65.00 3.300      3.2965      3.7747      0.4782389341
80.00 3.790      3.4882      3.1347      0.3535647457
75.00 3.660      3.4243      3.3480      0.0762968524
71.00 3.400      3.3732      3.5187      0.1455174622
66.00 2.740      3.3093      3.7321      0.4227853554
70.00 3.400      3.3604      3.5614      0.2009710408
72.00 3.550      3.3860      3.4760      0.0906638835
80.00 3.600      3.4882      3.1247      0.2535647457
71.78 3.490      3.3832      3.4854      0.1022636708

--- EXPERIMENT 2 (Attendance% -> GPA) ---
Att%  Actual GPA  Sklearn Pred  Manual Pred  Difference
85.00 8.000      3.2551      3.8105      0.5553983522
89.00 3.300      3.3185      3.6684      0.3499114963
100.00 3.790      3.4928      3.2776      0.2151773575
95.00 3.660      3.4136      3.4552      0.0416812124
91.00 3.400      3.3502      3.5974      0.2471686684
90.00 2.740      3.3343      3.6329      0.2985397823
96.00 3.400      3.4294      3.4197      0.0096905015
93.00 3.550      3.3019      3.5263      0.1444246404

```

```

# Statistical comparison metrics

# Calculate metrics for manual predictions
mse_cia_manual = mean_squared_error(y_test, y_pred_cia_manual)
r2_cia_manual = r2_score(y_test, y_pred_cia_manual)

mse_att_manual = mean_squared_error(y_att_test, y_pred_att_manual)
r2_att_manual = r2_score(y_att_test, y_pred_att_manual)

print("METRICS COMPARISON")

print("\n--- EXPERIMENT 1 (CIA% -> GPA) ---")
print(f'Metric: {<15} {<15} {<15} {<15}')
print(f'MSE: {<15} {<15.6f} {<15.6f}')
print(f'R^2 Score: {<15} {<15.6f} {<15.6f}')

print("\n--- EXPERIMENT 2 (Attendance% -> GPA) ---")
print(f'Metric: {<15} {<15} {<15} {<15}')
print(f'MSE: {<15} {<15.6f} {<15.6f}')
print(f'R^2 Score: {<15} {<15.6f} {<15.6f}')

```

Python

```

METRICS COMPARISON

--- EXPERIMENT 1 (CIA% -> GPA) ---
Metric      Scikit-learn  Manual OLS
MSE         3.017112    0.506235
R^2 Score   -0.549907    0.739931

--- EXPERIMENT 2 (Attendance% -> GPA) ---
Metric      Scikit-learn  Manual OLS
MSE         2.308560    1.885803
R^2 Score   -0.185981    0.031203

```

```

# Create dictionary with all model parameters
model_parameters = {
    'experiment_1': {
        'independent_variable': 'CIA Percentage',
        'dependent_variable': 'GPA',
        'slope': slope_cia,
        'intercept': intercept_cia,
        'r2_score': r2_cia,
        'mse': mse_cia
    },
    'experiment_2': {
        'independent_variable': 'Attendance Percentage',
        'dependent_variable': 'GPA',
        'slope': slope_att,
        'intercept': intercept_att,
        'r2_score': r2_att,
        'mse': mse_att
    }
}

# Save parameters to pickle file
file_loc = '../models/linear_regression_weights.pkl'
with open(file_loc, 'wb') as file:
    pickle.dump(model_parameters, file)

print(f'Parameters saved successfully to {file_loc}')
print("\nSaved parameters:")
print(model_parameters)

```

Python

```

Parameters saved successfully to ../models/linear_regression_weights.pkl

Saved parameters:
{'experiment_1': {'independent_variable': 'CIA Percentage', 'dependent_variable': 'GPA', 'slope': 0.01278314458206303, 'intercept': 2.4655814552535955, 'r2_score': -0.5499865790161718, 'mse': 3.0171124255047177}, 'experiment_2': {'independent_variable': 'Attendance Percentage', 'dependent_variable': 'GPA', 'slope': 0.0012345678901234567, 'intercept': 1.2345678901234567, 'r2_score': -0.185981031203, 'mse': 2.3085600000000005}}

```

```

# Load the saved parameters
with open(file_loc, 'rb') as file:
    loaded_params = pickle.load(file)

print("Parameters loaded successfully!")
print(loaded_params)

# Use loaded parameters for prediction
print("PREDICTIONS USING LOADED PARAMETERS")

# Example predictions
example_cia = np.array([75], [80], [85], [90])
example_att = np.array([90], [92], [95], [98])

# Get slopes and intercepts from loaded params
slope1 = loaded_params['experiment_1']['slope']
intercept1 = loaded_params['experiment_1']['intercept']
slope2 = loaded_params['experiment_2']['slope']
intercept2 = loaded_params['experiment_2']['intercept']

# Make predictions
pred_gpa_cia = slope1 * example_cia + intercept1
pred_gpa_att = slope2 * example_att + intercept2

print("\nPredictions using CIA:")
for cia, gpa in zip(example_cia.flatten(), pred_gpa_cia.flatten()):
    print(f"CIA: {cia:.0f} -> Predicted GPA: {gpa:.3f}")

print("\nPredictions using Attendance:")
for att, gpa in zip(example_att.flatten(), pred_gpa_att.flatten()):
    print(f"Attendance: {att:.0f} -> Predicted GPA: {gpa:.3f}")

```

Parameters loaded successfully!

```

loaded_params:
{'experiment_1': {'independent_variable': 'CIA Percentage', 'dependent_variable': 'GPA', 'slope': 0.01278314458206303, 'intercept': 2.465581452535955, 'r2_score': -0.549885796161718, 'mse': 3.6171124255847177}, 'experiment_2': {'independent_variable': 'Attendance', 'dependent_variable': 'GPA', 'slope': 0.01278314458206303, 'intercept': 2.465581452535955, 'r2_score': -0.549885796161718, 'mse': 3.6171124255847177}}

```

PREDICTIONS USING LOADED PARAMETERS

```

Predictions using CIA:
CIA: 75 -> Predicted GPA: 3.424
CIA: 80 -> Predicted GPA: 3.488
CIA: 85 -> Predicted GPA: 3.552
CIA: 90 -> Predicted GPA: 3.616

Predictions using Attendance:
Attendance: 90 -> Predicted GPA: 3.334
Attendance: 92 -> Predicted GPA: 3.366
Attendance: 95 -> Predicted GPA: 3.418
Attendance: 98 -> Predicted GPA: 3.461

```

Conclusion /inference

1. The Scikit-learn method and manual OLS calculation gave exactly the same slope and intercept values, confirming our manual calculations were correct
2. Among the two variables, CIA percentage had a higher R-squared value compared to attendance percentage, meaning CIA percentage is a better predictor of GPA
3. Both CIA and attendance percentage showed a positive relationship with GPA - when these percentages increase, GPA tends to increase slightly
4. The model parameters were successfully saved in a pickle file and we were able to load and use them for making new predictions

Lab 4

Lab Exercise IV:

Aim

- Implement KNN Classification: To build and apply a K-Nearest Neighbors model for classifying breast cancer as malignant or benign.
 - Data Splitting and Preparation: To correctly divide the dataset into training and testing sets to ensure robust model evaluation.
 - Heuristic K Selection: To understand and apply a basic method for choosing the optimal 'K' value for the KNN algorithm.
 - Comprehensive Performance Evaluation: To analyze the model's effectiveness using various classification metrics including Accuracy, Precision, Recall, F1-Score, and ROC-AUC.
 - Cross-Validation for Generalization: To employ cross-validation to assess the model's stability and generalization ability across different data subsets.
-

Evaluation Rubrics

Submission Guidelines

19. Make a copy of the lab manual template with your <name_reg:no_subject name>
20. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
21. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
22. Create a **Git repository** in your profile <ML lab-reg no>. Follow a different branch for each lab <Lab 1, Lab 2 ...> and push the code to Git.
23. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
24. Upload the PDF to Google Classroom before the deadline.

Code with Results

```
[1]
✓ 1s
# First, we need to import all the necessary libraries.
# 'sklearn' (scikit-learn) is a powerful library for machine learning.
# 'pandas' is good for handling data in tables.
# 'numpy' is for numerical operations.

import pandas as pd
import numpy as np
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, roc_curve
import matplotlib.pyplot as plt
import seaborn as sns
```

1. Loading the Breast Cancer Dataset

The Breast Cancer dataset is a classic dataset for binary classification. It contains features computed from digitized images of fine needle aspirate (FNA) of a breast mass. The goal is to classify whether the mass is malignant (cancerous) or benign (non-cancerous).

```
[2]
✓ 0s
# Load the dataset from scikit-learn.
# load_breast_cancer() returns a dictionary-like object.
cancer = load_breast_cancer()

# The data (features) are in 'cancer.data'.
# The target (what we want to predict) is in 'cancer.target'.
# The feature names (column headers) are in 'cancer.feature_names'.
# The target names (labels) are in 'cancer.target_names'.

X = cancer.data # Features
y = cancer.target # Target variable (0 for malignant, 1 for benign)

print("Shape of features (X):", X.shape)
print("Shape of target (y):", y.shape)
print("Feature names:", cancer.feature_names[5]) # Display first 5 feature names
print("Target names:", cancer.target_names) # 0: malignant, 1: benign

# Let's put the data into a pandas DataFrame for easier viewing
df = pd.DataFrame(X, columns=cancer.feature_names)
df['target'] = y

print("\nFirst 5 rows of the dataset:")
display(df.head())
```

Shape of features (X): (569, 30)
Shape of target (y): (569,)
Feature names: ['mean radius' 'mean texture' 'mean perimeter' 'mean area' 'mean smoothness']
Target names: ['malignant' 'benign']

First 5 rows of the dataset:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension	...	worst texture	worst perimeter	worst area	worst smoothness	worst compactness	worst concavity	worst concave points	worst symmetry	worst fractal dimension	target
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	0.07871	...	17.33	184.60	2019.0	0.1622	0.6656	0.7119	0.2654	0.4601	0.11890	0
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	0.05667	...	23.41	158.80	1956.0	0.1238	0.1866	0.2416	0.1860	0.2750	0.08902	0
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	0.05999	...	25.53	152.50	1709.0	0.1444	0.4245	0.4504	0.2430	0.3613	0.08758	0
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	0.00744	...	26.50	98.87	567.7	0.2098	0.8663	0.6869	0.2575	0.6638	0.17300	0
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	0.05883	...	16.67	152.20	1575.0	0.1374	0.2050	0.4000	0.1425	0.2364	0.07678	0

5 rows × 31 columns

3. Heuristic K Selection for KNN

In KNN, 'K' is the number of nearest neighbors considered when making a prediction. Choosing a good 'K' is important.

A simple heuristic (rule of thumb) for choosing 'K' is to take the square root of the total number of samples in your training data. If the result is an even number, it's often good practice to pick an odd number to avoid ties in classification (e.g., if K=2, and one neighbor is class A and one is class B, it's a tie). Let's calculate a 'K' value based on this heuristic.

```
[4]
✓ 0s
# Calculate K using the square root of the number of training samples.
# We will ensure K is an odd number.

num_training_samples = X_train.shape[0]
heuristic_k = int(np.sqrt(num_training_samples))

# Ensure K is odd
if heuristic_k % 2 == 0:
    heuristic_k = 1 # Subtract 1 to make it odd
    if heuristic_k == 0: # Handle case where original is 1, becomes 0
        heuristic_k = 1

print(f"Number of training samples: {num_training_samples}")
print(f"Heuristic K value (approx. sqrt): {heuristic_k}")

# For simplicity, we will use this heuristic k for our first model.
```

Number of training samples: 455
Heuristic K value (approx. sqrt): 21

4. Training the KNN Classifier

Now we will initialize our KNN classifier with the chosen 'K' value and train it using our training data.

- `n_neighbors` is the K value.
- `fit(X_train, y_train)` is where the model 'learns' by storing the training data.

```
[5]
✓ 0s
# Initialize the KNN classifier with our heuristic K.
knn_model = KNeighborsClassifier(n_neighbors=heuristic_k)

# Train the model using the training data.
knn_model.fit(X_train, y_train)

print(f"KNN model trained with K = {heuristic_k} neighbors.")
```

KNN model trained with K = 21 neighbors.

5. Making Predictions

After training, we use the model to make predictions on the unseen test data. This `y_pred` will be compared against the actual `y_test` to evaluate performance.

```
(5)
✓ Os
# Use the trained model to predict the classes for the test set.
y_pred = knn_model.predict(X_test)

# If you also need prediction probabilities (for ROC-AUC), use 'predict_proba'.
y_pred_proba = knn_model.predict_proba(X_test)[:, 1] # Probability of the positive class (class 1 - benign)

print("First 10 actual labels (y_test):", y_test[:10])
print("First 10 predicted labels (y_pred):", y_pred[:10])

First 10 actual labels (y_test): [1 0 0 1 1 0 0 0 1 1]
First 10 predicted labels (y_pred): [1 0 0 1 1 0 0 0 1 1]
```

6. Evaluating Model Performance: Classification Metrics

For classification tasks, we use specific metrics to understand how well our model is performing. Unlike regression metrics (like Mean Squared Error), classification metrics focus on how many predictions were correct or incorrect, and what kind of errors were made.

- **Accuracy:** The proportion of total correct predictions (both positive and negative) out of all predictions.
- **Precision:** Out of all predicted positive cases, how many were actually positive? (Minimizes false positives).
- **Recall (Sensitivity):** Out of all actual positive cases, how many did the model correctly identify? (Minimizes false negatives).
- **F1-Score:** The harmonic mean of Precision and Recall. It's a good measure when you need to balance both precision and recall, especially with imbalanced classes.

```
(7)
✓ Os
# Calculate various classification metrics

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")

--- Accuracy: 0.9649
Precision: 0.9467
Recall: 1.0000
F1-Score: 0.9726
```

6. Evaluating Model Performance: Classification Metrics

For classification tasks, we use specific metrics to understand how well our model is performing. Unlike regression metrics (like Mean Squared Error), classification metrics focus on how many predictions were correct or incorrect, and what kind of errors were made.

- **Accuracy:** The proportion of total correct predictions (both positive and negative) out of all predictions.
- **Precision:** Out of all predicted positive cases, how many were actually positive? (Minimizes false positives).
- **Recall (Sensitivity):** Out of all actual positive cases, how many did the model correctly identify? (Minimizes false negatives).
- **F1-Score:** The harmonic mean of Precision and Recall. It's a good measure when you need to balance both precision and recall, especially with imbalanced classes.

```
(7)
✓ Os
# Calculate various classification metrics

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)


print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")

--- Accuracy: 0.9649
Precision: 0.9467
Recall: 1.0000
F1-Score: 0.9726
```

7. ROC-AUC (Receiver Operating Characteristic - Area Under the Curve)

ROC-AUC is another powerful metric, especially for binary classification. It tells us how well the model distinguishes between positive and negative classes across all possible classification thresholds.

- The ROC curve plots the True Positive Rate (Recall) against the False Positive Rate at various threshold settings.
- The AUC (Area Under the Curve) measures the entire area underneath the ROC curve. An AUC of 1.0 means a perfect classifier, while an AUC of 0.5 means a classifier that performs no better than random guessing.



8. Cross-Validation

Cross-validation is a technique to get a more reliable estimate of model performance. Instead of a single train-test split, it divides the dataset into multiple 'folds' (e.g., 5 or 10).

Here's how K-fold cross-validation works:

- The dataset is split into K equal-sized folds.
- For each fold:
 - One fold is used as the **test set**.
 - The remaining K-1 folds are used as the **training set**.
- The model is trained and evaluated K times, each time with a different test fold.
- The results (e.g., accuracy scores) from each fold are averaged to get a more robust performance estimate.

This helps to reduce the impact of a particular data split and gives a better idea of how the model generalizes.

```
[9] In [ ]: # We will use 5-fold cross-validation.
# This means the data will be split into 5 parts, and the model will be trained and tested 5 times.

cv_scores = cross_val_score(knn_model, X, y, cv=5) # 'cv' specifies the number of folds

print("Cross-validation scores (accuracy for each fold):")
print(cv_scores)
print(f"Mean CV Accuracy: {np.mean(cv_scores):.4f}")
print(f"Standard Deviation of CV Accuracy: {np.std(cv_scores):.4f}")

... Cross-validation scores (accuracy for each fold):
[0.86842105 0.92105263 0.93859649 0.95614035 0.96460177]
Mean CV Accuracy: 0.9298
Standard Deviation of CV Accuracy: 0.0341
```

Comparison: Classification Metrics vs. Regression Metrics

This lab focused on **classification metrics** (Accuracy, Precision, Recall, F1-Score, ROC-AUC) which are used when you are predicting categories or labels (e.g., 'malignant' or 'benign'). These metrics tell us how many items were correctly assigned to their respective classes.

In contrast, **regression metrics** (like Mean Absolute Error (MAE), Mean Squared Error (MSE), Root Mean Squared Error (RMSE), R-squared) are used when you are predicting continuous numerical values (e.g., house prices, temperature). These metrics measure the difference between the predicted numerical value and the actual numerical value.

Key Difference:

- Classification:** Deals with discrete output (categories). Metrics measure correctness of class assignment.
- Regression:** Deals with continuous output (numbers). Metrics measure the error in numerical prediction.

Conclusion /inference

1. Dataset Overview: The Breast Cancer dataset contains 569 samples with 30 features, allowing for the classification of tumors into 'malignant' or 'benign'.
2. Heuristic K Value: Based on the training data size (455 samples), a heuristic value of $K = 21$ was selected for the KNN model.
3. Strong Model Performance: The KNN model demonstrated excellent performance on the test set, achieving:
 - a. Accuracy: ~96.5%
 - b. Recall: 1.0000 (meaning all actual benign cases were correctly identified).
 - c. F1-Score: 0.9726
4. High Discriminative Power: The ROC-AUC score of 0.9961 suggests that the model has a very strong ability to distinguish between malignant and benign cases.
5. Robust Generalization: The 5-fold cross-validation resulted in consistently high accuracy scores (with a mean of 0.9298 and a low standard deviation), indicating that the model is stable and likely to perform well on new, unseen data.